



# FINAL PROJECT

# CLOUD COMPUTING (BASIC)

Giulio Fantuzzi

# PROBLEM OVERVIEW



## PROJECT AIMS

- To identify and deploy a cloud-based file storage system
- To enable users to upload, download and delete files
- To address scalability, security and cost-efficiency
- To test the platform under stress



## WHY NEXTCLOUD?

- Simplicity in implementing the system
- Supported by an extensive documentation
- Being open-source, it benefits from a vast community of developers and users
- Comprehensive set of features beyond basic file storage

```
CloudBasic_FinalProject/
|
|   create_usr.sh
|
|   data/
|       ...files that will be created...
|
|   delete_usr.sh
|
|   docker-compose.yml
|
|   generate_files.sh
|
|   locust-tests/
|       locustfile.py
|
|   README.md
|
|   setup.sh
```

# USER AUTHENTICATION AND FILE OPERATIONS

Requirements about user authentication and authorization are guaranteed:



## User Management

- Nextcloud offers a built-in user management system
  - sign-up
  - login
  - logout
- User-friendly system, ensuring users to easily navigate the platform



## Role Based Access Control (RBAC)

- Different user roles:
  - Regular users
  - Admins
- Crucial to manage access levels in the system
- Admins have permissions to manage users
- Regular users have their private storage space



## Admin Management

- Admins manage users through an admin interface
  - Create users
  - Edit users
  - Delete users
  - Assigning/Change roles and permissions




## Private Storage Space

- Users are automatically assigned a private storage space upon account creation
- Default 512 MB storage quota-per-user
- Admins can adjust individual users' quotas through the web interface
- Global quota by modifying a *config.php* file

# ADDRESS SECURITY

- Nextcloud offers a comprehensive suite of security features that can be configured from the administrators interface
- All the options can be found in **Administration settings/Security/** and can be easily enabled
- Despite not being essential for the purposes of this exercises, they would be crucial for a real-world scenario



## Basic password policies

- Minimum password length
- User password history
- Number of days until a user password expire
- N. of login attempts before an account is blocked



## Server-side-encryption (SSE)

- Files are encrypted before being stored
- Files automatically encrypted upon upload
- Files automatically decrypted upon download
- Encryption process transparent to the end-users
- Some drawbacks in terms of performance, but fundamental to strengthen the overall security



## Enhanced password policies

- Forbid common passwords
- Enforce upper and lower case characters
- Enforce numeric/special characters
- Check passwords against the list of breached passwords from *haveibeenpwned.com*



## Two factor authentication (2FA)

- Adds an additional layer of protection
- It enforces users to provide two forms of identification: a password + a second factor
- It reduces unauthorized access to accounts
- Deliberately excluded from my analysis....

# COST-EFFICIENCY

- I referred to Nextcloud Server, which is the **free version** of Nextcloud
- It can be deployed on a dedicated server infrastructure, making it particularly suitable for small/medium businesses
- **Nextcloud Enterprise** for Larger organizations with more advanced needs and requiring professional support



Sources of costs

Hardware on which Nextcloud is hosted

Storage solutions used to store the data

- **Most intriguing challenge:** cost-efficiency in storage management
  - Organizations experiencing fluctuating data usage and uncertain business trajectory (startups)
  - Forecasting the necessary amount of storage often proves to be infeasible
- **Balanced approach** that incorporates both on-premises and cloud storage solutions:
  - On-premises storage to perform essential tasks independently of external providers
  - Cloud storage solution through a *pay-as-you-go* model according to demand

# ADDRESS SCALABILITY

- As the number of users and files increases, it becomes essential to implement a scalable file storage architecture
- Solution 1: **mixed strategy** combining on-premises and cloud storage solutions
- Solution 2: deploying a distributed environment that employs **horizontal partitioning of data**
  - Viable approach to enhance scalability
  - Data divided into smaller, more manageable chunks
  - Load-data distributed across multiple nodes and improving performance
- What about **NoSQL solutions** ?
  - In a distributed setup, NoSQL solutions like MongoDB may be intriguing
  - Scalability, flexibility, and performance for handling large volumes of unstructured data
- BUT feasibility not trivial at all
  - Nextcloud does not directly support NoSQL databases 😬 (it primarily works with SQL databases)
  - Transition to a NoSQL database is strictly dependent on some data requirements ⚠
- I equipped my Nextcloud instance with **Redis** a caching mechanism which allows to store frequently accessed data into memory, reducing the need to query the primary storage or database for every request




mongo DB



n8n.io

# DEPLOYMENT OF THE SYSTEM

The deployment of Nextcloud has been executed using **Docker** and **Docker Compose**



Docker images:

- **nextcloud**
  - To deploy the Nextcloud application
  - Configured with environment variables for database connection details, admin user credentials, and other settings
- **mariadb**
  - To deploy a MariaDB database server, providing the backend storage for Nextcloud
  - Configured with environment variables for setting up the root password, db name and user credentials
- **redis**
  - This image sets up a Redis server, which is utilized for caching purposes within the Nextcloud application
- **locustio/locust**
  - Open-source load testing tool, which is used to simulate user traffic and analyze the performance
  - Configured to run tests defined in a *locustfile.py* script targeting the Nextcloud instance.



# EXTERNAL CLOUD PROVIDER TO DEPLOY THE SYSTEM

For deploying a Nextcloud instance in a production environment, I would opt for **AWS (Amazon Web Services)** since it emerges as a compelling choice due to its scalability, reliability, and the wide range of services that it offers

- **Scalability and Performance**

- Highly scalable infrastructure, allowing to accommodate the growth as increase:
  - the user base
  - the storage needs
- Services as Amazon S3, which is highly scalable and can serve thousands of HTTP requests per second
- Crucial for handling varying loads and ensuring the Nextcloud instance remains responsive and available to users

- **Reliability**

- High reliability and availability
- Data centers located in various regions worldwide → Nextcloud instance remains accessible and functional

- **Cost-Effectiveness**

- *pay-as-you-go* model, which can lead to significant cost savings (already discussed)

- **Integration**

- AWS offers integration with a vast ecosystem of services and tools





# LOCUST TESTING



All locust tasks were defined in a Python script and executed on my MacBook Pro M2 laptop



Tests to assess the system's performance under varying stress conditions



Pool of 50 users

